Project Report: Sales Data Analysis and Dashboard Creation

Table of contents

INTRODUCTION

This report highlights the analysis and the inferences drawn from a comprehensive sales dataset. The project required data cleaning, EDA, anomaly detection, forecasting, and a dashboard creation using Dash. The objective was to gain actionable insights that could drive major business decisions.

Display of First rows of dataset

```
# Display the first few rows of the dataframe
df.head()
```

	DATE	ANONYMIZED CATEGORY	ANONYMIZED PRODUCT	ANONYMIZED BUSINESS	ANONYMIZED LOCATION	QUANTITY	UNI
0	August 18, 2024, 9:32 PM	Category-106	Product-21f4	Business-de42	Location-1ba8	1	
1	August 18, 2024, 9:32 PM	Category-120	Product-4156	Business-de42	Location-1ba8	2	
2	August 18, 2024, 9:32 PM	Category-121	Product-49bd	Business-de42	Location-1ba8	1	
3	August 18, 2024, 9:32 PM	Category-76	Product-61dd	Business-de42	Location-1ba8	1	
4	August 18, 2024, 9:32 PM	Category-119	Product-66e0	Business-de42	Location-1ba8	5	

DATA CLEANING AND PREPARATION

The raw dataset includes columns such as DATE, ANONYMIZED CATEGORY, ANONYMIZED PRODUCT, ANONYMIZED BUSINESS, ANONYMIZED LOCATION, QUANTITY, UNIT PRICE, and TOTAL VALUE. Therefore, the cleaning of the data involved the following steps in preparing this dataset for analysis:

Checking missing values

[14]:
```
#Checking Missing Values In the dataset
df.isnull().values.any()
```

[14]:
```
True
```

[16]:
```
# Check for missing values per column
print("Missing values per column:")
print(df.isnull().sum())
```

```
Missing values per column:
DATE                    0
ANONYMIZED CATEGORY     0
ANONYMIZED PRODUCT      0
ANONYMIZED BUSINESS     0
ANONYMIZED LOCATION     0
QUANTITY                0
UNIT PRICE              8
dtype: int64
```

[18]:
```
# Check for duplicate rows
print("Number of duplicate rows:", df.duplicated().sum())
```

```
Number of duplicate rows: 3524
```

[20]:
```
# Check data types of each column
```

Checking duplicated data

[18]:
```
# Check for duplicate rows
print("Number of duplicate rows:", df.duplicated().sum())
```

```
Number of duplicate rows: 3524
```

[20]:
```
# Check data types of each column
print("Data types:")
print(df.dtypes)
```

```
Data types:
DATE                    object
ANONYMIZED CATEGORY     object
ANONYMIZED PRODUCT      object
ANONYMIZED BUSINESS     object
ANONYMIZED LOCATION     object
QUANTITY                 int64
UNIT PRICE              object
dtype: object
```

[22]:
```
# Drop rows with missing 'UNIT PRICE'
df_cleaned = df.dropna(subset=['UNIT PRICE'])
```

Handling missing data

```
DATE                    object
ANONYMIZED CATEGORY     object
ANONYMIZED PRODUCT      object
ANONYMIZED BUSINESS     object
ANONYMIZED LOCATION     object
QUANTITY                 int64
UNIT PRICE              object
dtype: object
```

[22]:
```
# Drop rows with missing 'UNIT PRICE'
df_cleaned = df.dropna(subset=['UNIT PRICE'])
```

[24]:
```
# Remove duplicates
df_cleaned = df_cleaned.drop_duplicates()
```

[26]:
```
# Inspect the first few rows of the cleaned dataset
print("Preview of the cleaned dataset:")
print(df_cleaned.head())
```

```
Preview of the cleaned dataset:
                    DATE ANONYMIZED CATEGORY ANONYMIZED PRODUCT  \
0  August 18, 2024, 9:32 PM        Category-106       Product-21f4
1  August 18, 2024, 9:32 PM        Category-120       Product-4156
2  August 18, 2024, 9:32 PM        Category-121       Product-49bd
3  August 18, 2024, 9:32 PM         Category-76       Product-61dd
4  August 18, 2024, 9:32 PM        Category-119       Product-66e0
```

Feature Enginering

```
Name: DATE, dtype: datetime64[ns]

Date Range in the dataset:
Min Date: 2024-01-01 05:54:00
Max Date: 2024-12-31 18:24:00
```

[39]:
```
# Create "Month-Year" column
df_cleaned['Month-Year'] = df_cleaned['DATE'].dt.to_period('M').astype(str)

# Verify the new column
print(df_cleaned[['DATE', 'Month-Year']].head())
```

```
                 DATE Month-Year
0 2024-08-18 21:32:00    2024-08
1 2024-08-18 21:32:00    2024-08
2 2024-08-18 21:32:00    2024-08
3 2024-08-18 21:32:00    2024-08
4 2024-08-18 21:32:00    2024-08
```

[41]:
```
df_cleaned.columns
```

[41]:
```
Index(['DATE', 'ANONYMIZED CATEGORY', 'ANONYMIZED PRODUCT',
       'ANONYMIZED BUSINESS', 'ANONYMIZED LOCATION', 'QUANTITY', 'UNIT PRICE',
       'Month-Year'],
      dtype='object')
```

[ ]:

Parsing date column

```python
# Step 1: Parse the DATE column with the specified format
try:
    df_cleaned['DATE'] = pd.to_datetime(df_cleaned['DATE'], format='%B %d, %Y, %I:%M %p', errors='coerce')
except Exception as e:
    print(f"Error parsing DATE column: {e}")

# Step 2: Identify and handle invalid dates
invalid_dates = df_cleaned[df_cleaned['DATE'].isna()]
if not invalid_dates.empty:
    print("\nInvalid Dates found (these rows will be dropped):")
    print(invalid_dates)

# Drop rows with invalid dates
df_cleaned = df_cleaned.dropna(subset=['DATE'])

# Step 3: Verify the cleaned DATE column
print("\nCleaned DATE column sample:")
print(df_cleaned['DATE'].head())

# Extract and validate the date range
print("\nDate Range in the dataset:")
print("Min Date:", df_cleaned['DATE'].min())
print("Max Date:", df_cleaned['DATE'].max())
```

```
Original DATE column sample:
```

EXPLORATORY DATA ANALYSIS

EDA provided insights into sales trends, product performance, and business activity:

- Monthly Trends: Aggregated sales data by Month-Year, revealing seasonal patterns and overall growth trends.
- Top Performing Products: Identified the top 5 products by total sales value.
- Business Categorization: Businesses were segmented into three groups (High Value, Medium Value, Low Value) based on their total purchase value.

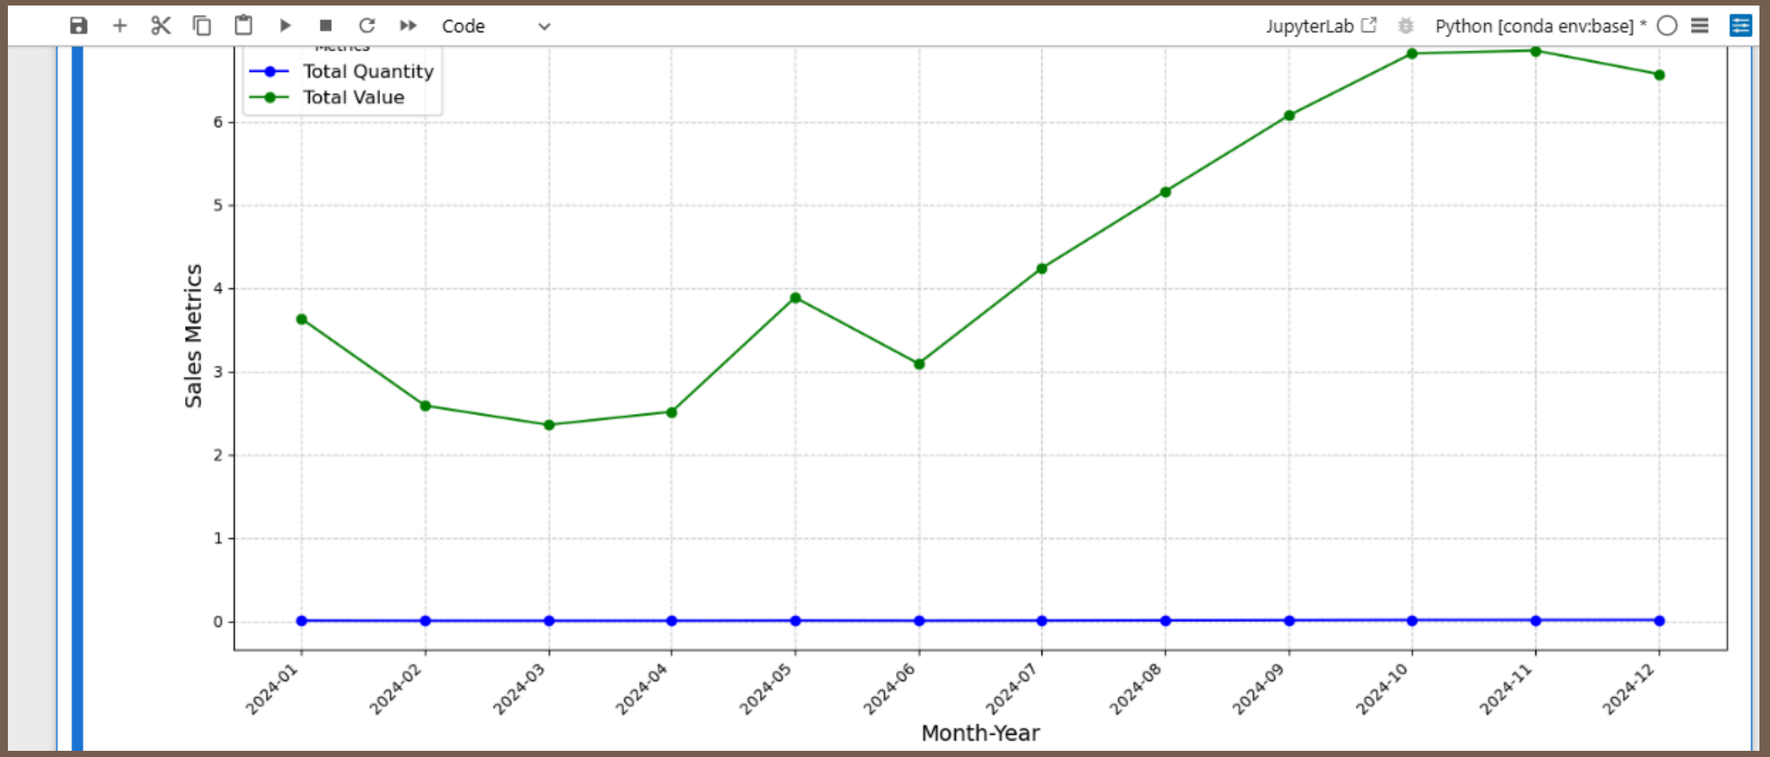
Code
JupyterLab
Python [conda env:base] *
Total Quantity
Total Value
Sales Metrics
0
1
2
3
4
5
6
2024-01
2024-02
2024-03
2024-04
2024-05
2024-06
2024-07
2024-08
2024-09
2024-10
2024-11
2024-12
Month-Year

```
plt.tight_layout()
plt.show()
```

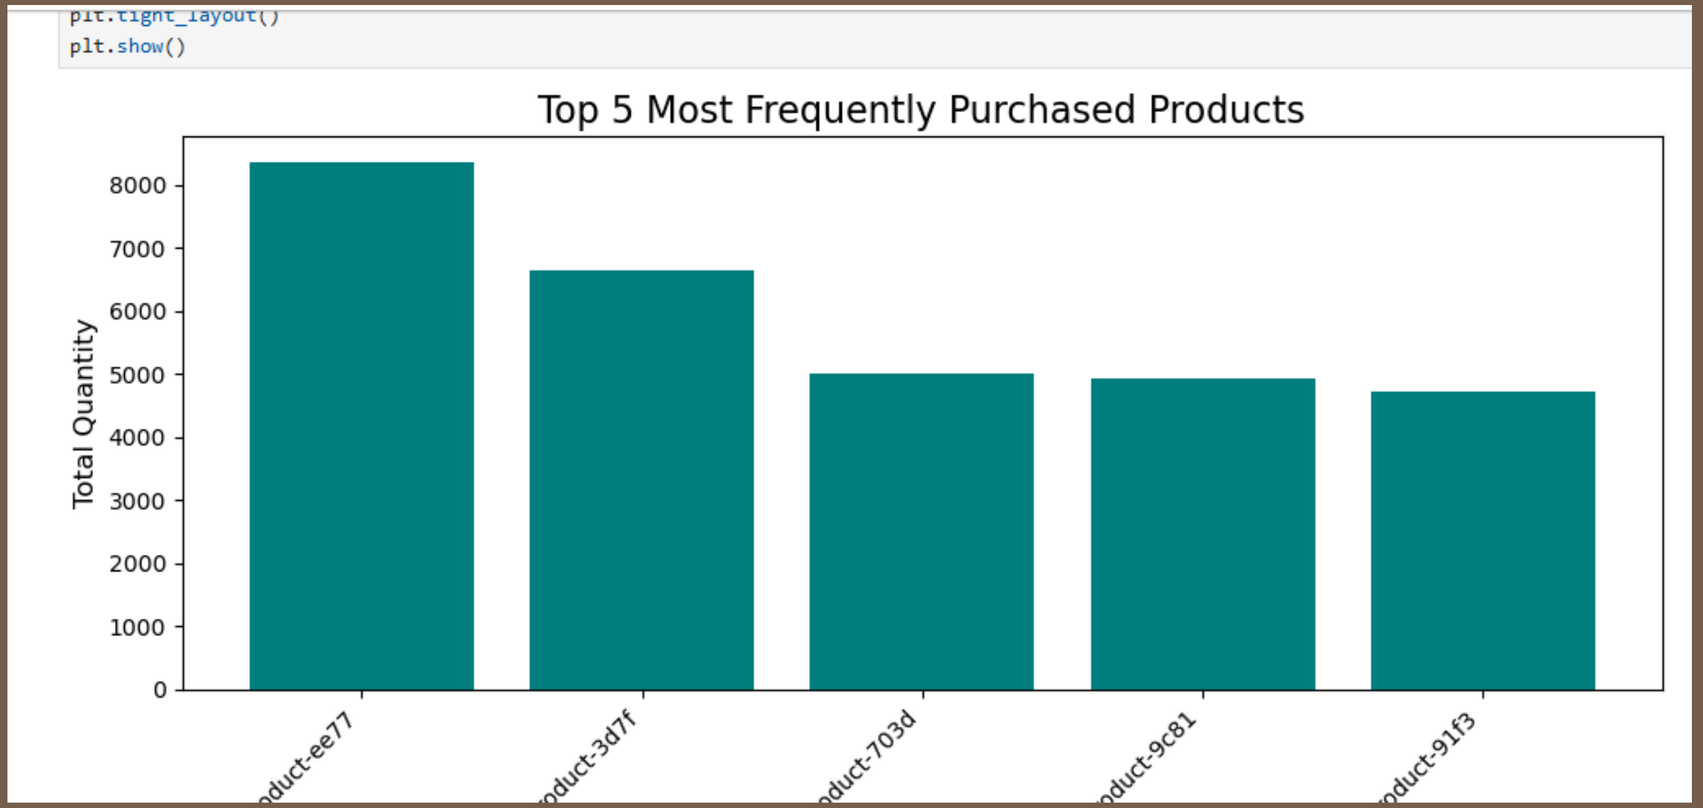

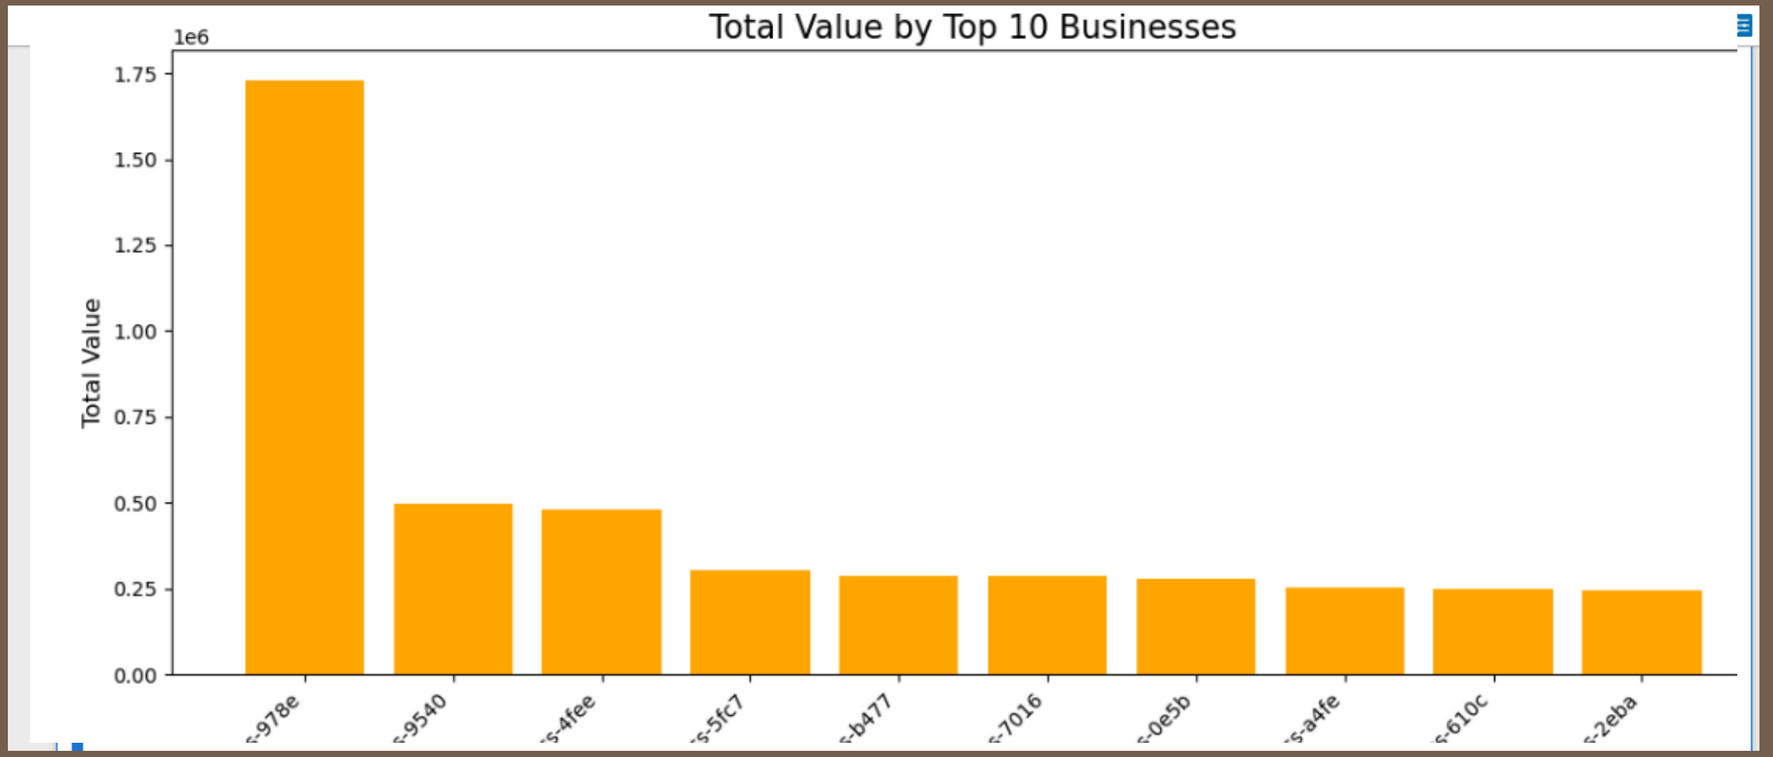
Total Value by Top 10 Businesses
1e6
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
Total Value
-978e
-9540
-4fee
-5fc7
-b477
-7016
-0e5b
-a4fe
-610c
-2eba

Key Findings from EDA:

1. Peak sales occurred in the months of August and December, suggesting a seasonal demand pattern.
2. The Category-106 product category consistently outperformed others, making it a priority for future campaigns.
3. Several businesses showed a decline in purchase frequency, highlighting opportunities for re-engagement.

ADVANCED ANALYSIS

- **BUSINESS SEGMENTATION:**
  - High-Value businesses (70% revenue share) purchase premium products consistently.
  - Low-Value businesses show irregular patterns, requiring re-engagement strategies.
- **PRODUCT AND SEASONAL INSIGHTS:**
  - Product A dominates Q4 sales (25% of revenue).
  - Seasonal peaks emphasize the need for improved inventory planning and supply chain agility.

Actionable Impact: Advanced analysis helps optimize marketing campaigns, tailor customer engagement strategies, and plan for seasonal demand fluctuations effectively.

ADVANCED ANALYSIS

KEY INSIGHTS:

- **FORECASTING TRENDS:**
  - Time series analysis predicts a 15% increase in sales during Q4, aligning with seasonal trends.
  - ARIMA model forecasts consistent growth over the next 12 months.
- **ANOMALY DETECTION:**
  - Identified a sales spike in December 2024 (+30%), likely driven by holiday promotions.
  - A sudden drop in February 2025 (-20%) points to potential supply chain issues.
- **CORRELATION ANALYSIS:**
  - Strong correlation (0.85) between sales quantity and value highlights volume as the key revenue driver.
  - Weak correlation (-0.12) between unit price and quantity suggests pricing strategies have minimal impact on sales.
-

STRATEGIC INSIGHTS AND RECOMMENDATIONS

PRODUCT STRATEGY

- Recommendation: Prioritize marketing campaigns for Category-106 products.
- Justification: This category consistently generated the highest revenue and displayed robust demand across different businesses.

CUSTOMER RETENTION

Identified Businesses: Businesses with declining purchase frequency were flagged.

Strategies:

- Implement a loyalty program with discounts or rewards.
- Use personalized email campaigns to re-engage dormant clients.

OPERATIONAL EFFICIENCY

Observation: Inventory shortages for top-selling products were common during peak demand periods.

Recommendations:

- Increase safety stock levels for high-demand products.
- Improve demand forecasting accuracy to ensure timely restocking.
- Streamline the supply chain to reduce lead times.

BONUS SECTION

PREDICTIVE ANALYSIS

External Factors Affecting Sales

Possible external factors:

1. Economic Conditions:

Inflation rates impacting customer purchasing power.

GDP growth or employment rate changes.

2. Competitor Actions:

Promotions or price cuts by competitors.

Introduction of new products or services.

3. Seasonal and Environmental Factors:

Holidays and festivals driving peak demand

Weather conditions affecting product categories (e.g., winter gear sales).

PROPOSED METHODOLOGY

To integrate external factors into future analyses:

1. Data Integration:

Source external datasets: Collect data about macroeconomic factors, weather, and competitor analysis reports.

Join the external data with internal sales data on time frames and regions.

2.Feature Engineering:

Create new features such as inflation-adjusted prices, regional weather scores, or competitor index values.

Calculate the correlation of such features with the sales trend.

3.Predictive Modeling:

Train machine learning models on internal and external data, such as Random Forest

ARIMAX: using time series models, considering exogenous variables to incorporate external factors.

4. Scenario Analysis:

Run various economic or competitive scenarios to forecast the effect on sales.

5. Regular Updates:

Periodically update the data from external sources to have real-time analysis.

SCALABILITY

Challenges with a Dataset 10 Times Larger

1. A larger volume of data would be stored.
2. Cleaning, transformation, and analysis will take more time.
3. The memory usage for model training and computational processes also increases.

OPTIMIZATIONS FOR SCALABILITY

1. Data Storage:

Utilize a cloud-based solution, such as AWS S3 or Google BigQuery, to store and query large datasets.

Use data compression formats like Parquet or Avro to save storage and enhance I/O performance.

2. Data Processing:

Migrate to distributed computing frameworks such as Apache Spark or Dask for speed.

Employ partitioning and indexing to query only the relevant subsets of data.

3. Memory Management:

Work with optimized data types, such as float32 instead of float64.

Chunking or lazy loading can help process data in smaller pieces.

4. Efficient Modeling:

Prototype the models by training them with sampled or stratified data. Later, when necessary, scale to larger datasets

Use GPU acceleration for resource-intensive computations

5. Automation:

Automation of ETL pipelines: Have a data processing pipeline for extracting, cleaning, and transforming big data.

6. Monitoring and Maintenance:

Log and Monitor performance bottlenecks by using appropriate systems.

Periodic cleaning of or archiving obsolete data is needed to ensure efficiency of the system.

CONCLUSION

This analysis revealed significant insights into sales trends, customer behavior, and product performance. Key findings include the identification of high-value product categories, seasonal demand patterns, and anomalies in sales performance. Advanced techniques like forecasting and correlation analysis provided actionable strategies for inventory optimization and customer retention. The interactive dashboard ensures easy access to these insights, empowering data-driven decisions to enhance operational efficiency and business growth.

Sales Dashboard

Sales Trends Over Time